Universal Template Parameters

Contents

1 Introduction

We propose a universal template parameter (UTP) and associated cleanups to make it useful for template metaprogramming. This will allow for a generic `apply` and other higher-order template metafunctions, including certain type traits.

2 Change Log

2.1 R2 -> R3

    — Complete paper rewrite to make it more obvious what is proposed and what has been considered.

— Proposed examples have been expanded.
— Added universal aliases
— Merged variable template template parameters paper into this one
— Merged the concept template parameters paper into this one
— Respecified UTPs as dependent names and continued cleanup started with "down with `typename`"

2.2   R1 -> R2

Having found overwhelming support for the feature in EWGI, and a concern about contra- and co-variance of the `template auto` parameters, we include discussion of these topics in the paper, together with the decision and comparison tables that underpin the decision.

The decision was to explore both and put it to a vote.

Added list of questions for EWGI.

2.3   R0 -> R1

— Greatly expanded the number of examples based on feedback from the BSI panel.
— Clarified that we are proposing eager checking

3   Related work

— [P0634R3] [accepted] got rid of `typename` where it was obviously redundant; we hope to get rid of it in even more places (though not all).
— [P0945R0] [discarded] explored universal aliases; this proved unworkable, and the minutes of discussion informed this paper.
— [P2601R1] explores dropping empty `<>`. This might take up syntactic space that we need, but we haven't explored whether it does yet due to lack of time.
— [P0522R0] explores related partial ordering around partial template specialization; this paper to author's knowledge does not interact with that paper.
— [P2008R0] proposes variable-template template-parameters

4   Motivation and Examples

This paper unifies the model of template parameters with the model for dependent tokens (types, values, templates, and at some point hopefully concepts).

This model is not uniform in C++23, because it lacks a way to treat all of the above uniformly, ironically denying the ability of generic code to treat itself generically.

Note: `template auto` is a placeholder syntax for such a parameter, albeit not a bad one; see the spelling discussion section later in the paper.

Consider the following examples this paper aims to enable.

4.1   Checking whether a type is a specialization of a given template

Main discussion of feature in [P2098R0] by Walter Brown and Bob Steagall.

When writing template libraries, it is useful to check whether a given type is a specialization of a given template. When our templates mix types and NTTPs, this trait is currently impossible to write in general. However, with a universal template parameter (UTP), we can write a concept for that easily as follows.

```cpp
// is_specialization_of
template <typename T, template <template auto...> typename Type>
constexpr bool is_specialization_of_v = false;
```

```cpp
template <template auto... Params, template <template auto...> typename Type>
constexpr bool is_specialization_of_v<Type<Params...>, Type> = true;

template <typename T, template <template auto...> typename Type>
concept specialization_of = is_specialization_of_v<T, Type>;
```

This enables constraining to specific class templates:

```cpp
// example from a units library
template <auto N, auto D>
struct ratio {
  static constexpr decltype(N) n = N;
  static constexpr decltype(D) d = D;
};

template <specialization_of<ratio> R1, specialization_of<ratio> R2>
using ratio_mul = simplify<ratio<
    R1::n * R2::n,
    R1::d * R2::d
  >>;

// std::array<class, size_t>
static_assert(specialization_of<std::array<int, 4>, std::array>);
// std::vector<class, class, class>
static_assert(specialization_of<std::vector<int>, std::vector>);
```

4.2 `apply1`

A contrived-for-simplicity example which avoids the complexity of variadics:

```cpp
template <template <template auto> typename F, template auto Arg>
using apply1 = F<Arg>;

template <typename X> struct takes_type {};
template <auto X>     struct takes_value {};
template <template <template auto...> typename X> struct takes_template {};

using r1 = apply1<takes_type, int>; // takes_type<int>
using r2 = apply1<takes_value, 3>;  // ok, takes_value<3>
using r3 = apply1<takes_template, takes_template>; // takes_template<takes_template>
```

4.3 Full `apply` metafunction

The non-contrived example is the `apply` metafunction, achievable like so:

```cpp
template <template <template auto...> typename F, template auto... Args>
using apply = F<Args...>; // easy peasy!

using r1 = apply<std::array, int, 3>; // ok, std::array<int, 3>
// ok, r2 is std::vector<int, std::pmr::allocator>
using r2 = apply<std::vector, int, std::pmr::allocator>;
```

In C++23, this is impossible to do; the various metaprogramming libraries get around that by boxing, or by only supporting type-taking metafunctions.

4.4 New Traits (also an example)

We should be able to define new kind-traits based on this feature. Note the partial specialization mechanism should work as one would expect.

Note: We define variable templates first because they compile faster.

```cpp
template <template auto> constexpr bool is_typename_v    = false;
template <typename T>    constexpr bool is_typename_v<T> = true;
template <template auto> constexpr bool is_value_v       = false;
template <auto V>        constexpr bool is_value_v<V>    = true;
template <template auto> constexpr bool is_template_v    = false;
template <template <template auto...> typename A>
constexpr bool is_template_v<A>                          = true;

// if we gain variable-template template-parameters
template <template auto>
constexpr bool is_variable_template_v                    = false;
template <template <template auto...> auto A>
constexpr bool is_variable_template_v<A>                 = true;

// if we gain concept template-parameters
template <template auto> constexpr bool is_concept_v     = false;
template <template <template auto...> concept A>
constexpr bool is_concept_v<A>                           = true;

// and their associated types:
template <template auto X> struct is_typename : std::bool_constant<is_typename_v<X>> {};
template <template auto X> struct is_value    : std::bool_constant<is_value_v<X>> {};
template <template auto X> struct is_template : std::bool_constant<is_template_v<X>> {};
template <template auto X> struct is_variable_template
    : std::bool_constant<is_variable_template_v<X>> {};
template <template auto X> struct is_concept
    : std::bool_constant<is_concept_v<X>> {};
```

4.5 A variable-to-type adaptor that exposes `::result`

A `box` is an important way of bridging to type-based metaprogramming, so we need to define it, also for the purposes of further examples.

Note: we are not proposing `box` for the standard, but a paper probably should.

Note: `box` would be less necessary if we had variable-template template-parameters.

```cpp
template <template auto> struct box; // impossible to define body

template <auto X>
struct box<X> { static constexpr decltype(X) result = X; };

template <typename X>
struct box<X> { using result = X; };

template <template <template auto...> typename X>
struct box<X> {
  template <template auto... Args>
  using result = X<Args...>;
};
```

```
// if we get variable-template template-parameters
template <template <template auto...> auto X>
struct box<X> {
  template <template auto... Args>
  static constexpr decltype(X<Args...>) result = X<Args...>;
};

// if we get concept template-parameters
template <template <template auto...> concept X>
struct box<X> {
  template <template auto... Args>
  concept result = X<Args...>;
};
```

4.6 `map_reduce` :: Total Example

We believe all the required features are used in the `map_reduce` example, where we have to map metafunction results into a `::result` member (we don't have variable template parameters yet).

```
template <template <template auto> typename Map,
          template <template auto...> typename Reduce,
          template auto... Args>
using map_reduce = Reduce<Map<Args>::result...>;
```

Note: notice that the above is a type-alias template. `box` allows us to return anything because it's a type.

Note: notice we expect the metafunction `Map` to return a `box`.

Let's count the number of types in the argument list, for instance:

```
template <int... xs> using sum = box<(0 + ... + xs)>;
template <template auto X> using boxed_is_typename = box<is_typename_v<X>>;
static_assert(2 == map_reduce<boxed_is_typename, sum, int, 1, long, std::vector>::result);
```

With variable-template template-parameters, we don't need to box `is_typename_v`, and we can define `map_reduce` slightly better:

```
template <template <template auto> typename Map,
          template <template auto...> template auto Reduce,
          template auto... Args>
using map_reduce_better = Reduce<Map<Args>...>;
```

and we can use it more comfortably as well:

```
static_assert(2 == map_reduce_better<is_typename_v, sum, int, 1, long, std::vector>::result);
```

Note: alas, we still need `sum` to be a box. We won't ever have universal aliases, but we *can* have universal dependent expressions.

Note: the kind of `Reduce`.

5 Mechanism

This chapter describes how universal template parameters (UTPs) work. Effectively, a UTP acts like a dependent name. This paper cleans up dependent expressions so that they are more useful everywhere.

5.1 Specializing class templates on parameter kind

UTPs introduce a similar generalizations as the `auto` universal NTTP did; in order to make it possible to pattern-match on the parameter, class templates need to be able to be specialized on the kind of parameter as well:

```cpp
template <template auto> struct X;

template <typename T>
struct X<T> {
  // T is a type
  using type = T;
};

template <auto val>
struct X<val> {
    using type = decltype(val);
};

template <template <typename> typename F>
struct X<F> {
    // F is a unary metafunction
    template <typename T>
    using type = F<T>;
};

template<template auto U> user()
{
    using type = X<U>;
}
```

This basic mechanism allows the utterance of a UTP only in:

— *template-parameter-list* as the declaration of such a parameter.
— *template-argument-list* as the usage of such a parameter. It can bind only to a template parameter declared `template auto`.

This allows building enough traits to connect the new feature to the rest of the language with library facilities, and rounds out template parameters as just another form of a compile-time parameter. However, it has the same kind of limitations we saw before `if constexpr` was introduced: Function templates often had to rely on helper structs to do simple things.

5.2 Allowing UTPs in code

To make this feature useful, UTPs must be usable in code. To make parsing possible we have to disambiguate UTPs according to dependent name rules.

We have a version of this problem already with overloaded template functions and dependent names:

```cpp
template<typename T> int f() { return sizeof(T); } // f#1
template<auto V> int f() { return V; } // f#2

template<typename T> int caller()
{
    std::vector<T::name> x; // error in C++23, for little reason.
    return f<T::name>();    // "down with typename"++
    return f<T::name*>();    // Error: T::name* is syntactically not a constant expression
}
```

In C++23, the `f#2` is called if `T::name` is a value; but if a type, it is an error, instead of dispatching to `f#1`.

Note: No C++23 code is broken by this change as `caller` is only callable when `T::name` resolves to a value in C++23.

With UTPs this problem is more articulated:

```cpp
template<template auto U> int caller2()
{
    return f<U>();
}
```

This paper proposes we keep *parsing* dependent expressions as-if they are values (C++23 rules) (unless disambiguated with `typename` and `template`), but we actually defer the kind check until instantiation time.

5.2.1 Deferring kind checking to instantiation

Dependent expressions are still parsed as *constant-expression*, but kind-checks are always deferred to substitution time.

UTPs are just dependent expressions.

```cpp
template<template auto U> int caller3()
{
    auto u = f<U>();          // Instantiates for values and types
    auto v = f<U*>();         // Error: U* is not syntactically a constant expression
    auto v = f<typename U*>(); // OK: U* is disambiguated
    using type = X<U>;        // Instantiates for all kinds.

    // Can fail instantiation if argument kind is wrong.
    auto v = U;               // Parses U as a value.
    using t = U*;             // Parses U as a type thanks to "down with typename".

    // Error during parse
    U x;                      // Error: U is parsed as a value
    template<typename T> using tpl = U<T>;  // Error: U is parsed as a value.
    // Disambiguation useful outside of template argument:
    template<typename T> using tpl = template U<T>; // OK if U is bound to a class template
    typename U x;
    typename template U<int> i;  // Note: New double disambiguation
    auto vi = template U<int>;   // OK if U bound to a variable-template taking a typename
}
```

The promotion of dependent names to universal template parameters also applies to template argument packs and enables utility `struct`s to perform transformations on packs of template parameters involving mixes of types and NTTPs:

```cpp
template<typename T>
struct unwrap
{
  using result = T;
};

template<typename T, T t>
struct unwrap<std::integral_constant<T, t>>
{
```

```
    static constexpr T result = t;
};
```

Using the above, this is valid:

```
template <template <template auto...> typename T, typename... Params>
using apply_unwrap = T<unwrap<Params>::result...>;


apply_unwrap<std::array, int, std::integral_constant<std::size_t, 5>> arr;
```

6 Universal aliases

A universal alias is a name given to a dependent name or universal template parameter. The universal alias is in itself a dependent name, just like a UTP.

The grammar for a universal alias is simply:

universal-alias: **template auto** *identifier* = *template-argument* ;

The box example is now trivial, but also unnecessary.

```
template<template auto U> struct box {
    template auto result = U;
};
```

With this definition the `map_reduce` example can be further refined:

```
template<template<template auto> template auto Map,
         template<template auto...> template auto Reduce,
         template auto... Args>
template auto map_reduce_best = Reduce<Map<Args>...>;
```

and we can use it more comfortably as well:

```
template<int... xs> constexpr int sum = (0 + ... + xs);
static_assert(2 == map_reduce_best<is_typename_v, sum, int, 1, long, std::vector>);
```

Note: `Map` is also declared as a `template auto` metafunction, meaning that it could be a class template or as in the case of `sum` a variable template.

Note: Now we no longer need `sum` to be a box. `map_reduce_best` is a universal alias template. We still need to disambiguate these when used. In this example, however, as the result of `sum` is a value, we don't have to disambiguate inside the `static_assert`.

6.1 Properties of universal aliases

6.2 Universal aliases are purely compile-time entities

— if values, their initializer is `consteval`
— if types, that is trivially true.

In other words, a universal alias behaves like a UTP, but one that is introduced by a declaration, as opposed to introduced as a template parameter.

The grammar production with *template-argument* as the initializer is correct, with one of its alternatives being *constant-expression.*

Note: in a previous iteration of [P0945R0], the initializer of value kind did not have to be `constexpr`, and the universal alias was just another name for a variable or function.

6.3 Universal aliases are always treated as dependent names

This is regardless of whether their initializer is dependent or not. This is to allow changing whether the initializer is dependent without impact to parsing.

7 Automatic disambiguation when :: is applied

To ensure that UTPs work exactly as dependent names we need to clean up how dependent names work. This can be done without breaking backwards compatibility, it just makes a few more cases valid, by generalizing a rule.

The lengthy description below sums up to this TL;DR style conclusion:

Let :: automatically disambiguate its left hand side as a type, and let `typename` disambiguate the last dependent name as a type.

This is enough to make everything work consistently and unsurprisingly for both dependent names and UTPs.

7.1 Digging into the nested disambiguation jungle

The situation for disambiguating nested dependent names in C++23 is as follows. Consider this example code:

```
template<typename T> struct S {
    typename T::type1::type2 v1;      // OK
    int x1 = T::type1::value1;        // Error
};
```

Only types (and namespaces and enumerations) can have named members accessible with operator ::.

One would imagine, therefore, that `T::type1` could be automatically disambiguated as one of those so that `::value1` could be applied.

This is not the case:

The `typename` in the declaration of `v1` disambiguates *both* `type1` and `type2` to be `typename`s. The error in the initializer of `x1` is due to `type1` not being treated as a type.

This leaves us with the unfortunate situation where `value1` can't be used as a value. A workaround is to extract `T::type1` into a type alias, and apply the `::value1` portion to the alias; then `value1` can be treated as a value by *not* prefixing `typename`:

```
template<typename T> struct S {
    using type1 = T::type1;           // Intermediate name
    int x2 = type1::value1;           // Access dependent name value1 as a value.
};
```

If we take a look at dependent names which are subordinate to a dependent name disambiguated as a template we have a similar situation, but not exactly equivalent:

```
template<typename T> struct S {
    typename T::template tpl<int, 2> v2;
    typename T::template tpl<int, 2>::type2 v3;
    int x3 = T::template tpl<int, 2>::value1;  // Works, but why?
};
```

For `v2` the leading `typename` disambiguates the template instance as being a type. This is required as the template instance would otherwise be assumed to be an instance of a variable template.

For `v3`, `typename` disambiguates *both* the type-instance and its member `type2` as types, consistent with the `v1` declaration.

How `x3` can be initialized when `x1` can't is hard to understand, maybe the reasons can be traced back to a time when variable templates were not a thing, and thus `tpl<int, 2>` must be a type, which allows `::value1` to be applied. This is not consistent with the non-template case `x1`.

We can also reverse the order between the template and the `typename`, so that `T` must contain a type containing a template. This template can be a class template or variable template as above.

```
template<typename T> struct S {
    typename T::type1::template tpl<int, 3> v4;
    int x4 = T::type1::template tpl<int, 3>;  // Works, but why?
};
```

Just as for the opposite order of template and type this works, the leading `typename` again disambiguates both `type1` and the template instance as types when declaring `v4`, while the `x4` initialization remains inconsistent with `x1`.

Clearly, if we want UTPs to work like dependent names we don't want dependent names to behave this strangely. Below all the cases above are listed inside one class template taking a universal template parameter `U`.

```
template<template auto U> struct S {
    // Please note every line is independent :)

    int x = U;      // U is treated as a value
    typename U v5; // U can be disambiguated to typename.

    // U can be disambiguated as variable template.
    int x5 = template U<int, 3>;
    // U can be disambiguated to a class template
    typename template U<int, 3> v5;

    // U can be a type containing a value or subtype
    int x6 = U::value1;
    typename U::type2 v6;

    // U can be disambiguated as a template.
    // The template instance is automatically disambiguated
    // as a type if followed by ::
    int x7 = template U<int, 3>::value1;
    typename template U<int, 3>::type2 v7;

    // Finally U can be disambiguated as a type containing a template, which is
    done by the trailing :: preceeding the template keyword.
    int x8 = U::template tpl<int, 3>;
    typename U::template tpl<int, 3> v8;
};
```

Deciding that `U::value1` is a value makes UTPs work smoother than nested dependent names do in C++23; in C++23 `int x1 = T::type1::value1;` is invalid, whereas this paper defines to be valid, as long as `T::type1` ends up being a type at instantiation time, and is parsed as a type at parse time.

To rectify this we should change the rule for `typename` disambiguation to only affect the last of a set of nested dependent names, while a `::` automatically disambiguates a dependent name to its left as a type.

This does not change the meaning of the `v1` declaration above but makes the `x1` initializer a value as `type1` is disambiguated by the trailing `::` and `value1` is a value as there is no leading `typename`. The `x3` and `x4` initializers behave the same as today, but with a more understandable rationale.

In the template cases `x5`, `x7`, `v5` and `v7`, we must explicitly specify `template` to correctly parse `<` as the *template-*

head introducer. This is novel (`template` could only appear after `::` in C++23), but utterly unsurprising.

```
int x5 = T::template name<int, 3>
typename T::template name<int, 3> v5;

int x7 = T::template name<int, 3>::value1;
typename T::template name<int, 3>::type2 v7;
```

Note: the proposed syntax is close to explicit template instantiation.

```
template SomeClassTemplate<int, 3>;
```

This is an explicit instantiation of a class template `SomeClassTemplate`.

Variable declarations are allowed in namespace scope, which in the presence of universal alias templates, comes close to the proposed syntax:

```
template <typename auto ... >
template auto U = int; // for the purposes of this example

typename template U<int, 3> x;
```

8   Variable and Concepts template template parameters.

Initially proposed in [P2008R0], variable template template parameters would allow passing variable templates, such as value type traits as template parameter.

Because concepts almost act like variable template of type `bool`, while offering more capabilities, we should allow concept template template parameters.

This allows to express important ideas such as `range_of<std::integral>` or `tuple_like<std::regular>`

```
template <typename R, template<typename> concept C>
concept range_of =
        ranges::input_range<R>
        && remove_cvref_t<ranges::range_reference_t<R>>>;
```

The syntax for these new entities is straight forward (and live little room for invention)

```
template <auto N> // Variable template parameter
template <template <> typename> // Type template template parameter
template <template <> auto>      // Variable template template parameter
template <template <> concept>  // Concept template template parameter
```

We can not have `template <concept>` as a concept is by definition a template.

Variable template of a specific type (`template <template <> int>`) are also impossible, as the type of a template variable can vary across specializations, and there is no real way to restrict specializations from doing that.

In addition of the "sequence of entity satisfying some concept" use cases (like `range_of<integral>`), we can use this feature to better express complex, repetitive concepts. Barry Revzin offered a number of extremely compelling such use cases in this blog post. For example, we can describe more succintly and expressively the many "indirect" iterator concepts used to support projections.

```
template<class F, class I, template <typename> concept direct>
concept indirect =
    indirectly_readable<I> &&
    copy_constructible<F> &&
    direct<F&, iter_value_t<I>&> &&
```

```
    direct<F&, iter_reference_t<I>> &&
    direct<F&, iter_common_reference_t<I>> &&
    common_reference_with<
      invoke_result_t<F&, iter_value_t<I>&>,
      invoke_result_t<F&, iter_reference_t<I>>>
;

template<class F, class I>
concept indirectly_unary_invocable =
    indirect<F, I, invocable>;

template<class F, class I>
concept indirectly_regular_unary_invocable =
    indirect<F, I, regular_invocable>;
```

Universal template parameters and variable/concept template parameters are orthogonal features, in that they could be standardized separately.

However, when combined, these features are even more expressive. Notably, we can support both `range_of<int>` and `range_of<std::integral>`, by combining both features. IE:

```
template <typename R, template auto T> // Primary universal template
constexpr bool is_range_of = delete;

template <typename R, template <typename> concept C> // Specialization for concepts
constexpr bool is_range_of<R, C> = C<R>;

template <typename R, typename T> // Specialization for concrete types
constexpr bool is_range_of<R,T> = std::is_same_v<R, T>;

template <typename R, template auto T>
concept range_of = is_range_of<std::remove_cvref_t<std::ranges::range_reference_t<R>>, T>;

// We can now constrain a range to a specific type
static_assert(range_of<std::string, char>);
// Or a concept
static_assert(range_of<std::string, std::integral>);
```

In addition of the compelling synergy, it would be hard to ensure these feature integrate well with one another, given the interactions. In particular, it means that:

— Universal parameters can be used as parameter to concept/variable template template parameters.

```
// A pack of variable template template parameter
// parametrized on universal template parameters
template <template<template auto...> auto... V>
```

`template auto` and `auto` are not ambiguous here, but it might be confusing for the reader.

A UTP can then be:

— A type
— A value
— A type template
— A variable template
— A concept

We need to ensure than the product of these interactions is coherent, especially for dependant expressions (Even

if neither variable or concept template template parameters should not require new disambiguators, `template` is enough to disambiguate all types of template template parameters). The best way to ensure that coherence is to evolve both feature at the same time, in the same paper.

9 Clarifying examples

9.1 Single parameter examples

```
template <int>                          struct takes_int {};
template <typename T>                   using  takes_type = T;
template <template auto>                struct takes_anything {};
template <template <typename> typename F> struct takes_metafunc {};

template <template <template auto> typename F, template auto Arg>
struct fwd {
  using type = F<Arg>; // ok, passed to template auto parameter
}; // ok, correct definition

void f() {
  fwd<takes_int, 1>{};      // ok; type = takes_int<1>
  fwd<takes_int, int>{};    // error, takes_int<int> invalid
  fwd<takes_type, int>{}; // ok; type = takes_type<int>
  fwd<takes_anything, int>{}; // ok; type = takes_anything<int>
  fwd<takes_anything, 1>{};   // ok; type = takes_anything<1>
  fwd<takes_metafunc, takes_int>{}; // ok; type = takes_metafunc<takes_int>
  fwd<takes_metafunc, takes_metafunc>{}; // error (1)
}
```

(1): `takes_metafunc` is not a metafunction on a *type*, so `takes_metafunc<takes_metafunc>` is invalid (true as of *C++98*).

9.2 Pack Expansion Example

It is interesting to think about what happens if one expands a non-homogeneous pack of these. The result should not be surprising:

```
template <template auto X, template auto Y>
struct is_same : std::false_type {};
template <template auto V>
struct is_same<V, V> : std::true_type {};

template <template auto V, template auto ... Args>
struct count : std::integral_constant<
    size_t,
    (is_same<V, Args>::value + ...)> {};

// ok, ints = 2:
constexpr size_t ints = count<int, 1, 2, int, is_same, int>::value;
```

9.3 Example of parsing ambiguity (if late check)

The reason for eager checking is that not doing that could introduce parsing ambiguities. Example courtesy of [P0945R0], and adapted:

```
template <template auto A>
struct X {
  void f() { A * a; }
};
```

The issue is that we do not know how to parse `f()`, since `A * a;` is either a declaration or a multiplication.

If we treated `A` as just a dependent name, this always parses as a multiplication. To treat A as a typename and a as a variable being declared the A has to be disambiguated as any dependent name.

Original example from [P0945R0]:

```
template <typename T> struct X {
  using A = T::something; // P0945R0 proposed universal alias
  void f() { A * a; }
};
```

10   Example Applications

This feature is very much needed in very many places. This section lists examples of usage.

10.1   Enabling higher order metafunctions

This was the introductory example. Please refer to the [Proposed Solution].

Further example: `curry`:

```
template <template <template auto...> typename F,
          template auto ... Args1>
struct curry {
  template <template auto... Args2>
  using func = F<Args1..., Args2...>;
};
```

10.2   Making dependent `static_assert(false)` work

Dependent static assert idea is described in [P1936R0] and [P1830R1]. In the former the author writes:

> Another parallel paper [P1830R1] that tries to solve this problem on the library level is submitted. Unfortunately, **it cannot fulfill all use-case since it is hard to impossible to support all combinations of template template-parameters in the dependent scope**.

The above papers are rendered superfluous with the introduction of this feature. Observe:

```
// stdlib
template <bool value, template auto... Args>
constexpr bool dependent_bool = value;
template <template auto... Args>
constexpr bool dependent_false = dependent_bool<false, Args...>;

// user code
template <template <typename> typename Arg>
struct my_struct {
  // no type template parameter available to make a dependent context
  static_assert(dependent_false<Arg>, "forbidden specialization.");
};
```

However, a language change, such as proposed by [P2593R0] would still be beneficial.

10.3 Universal alias as a library class.

While this paper does not try to relitigate Richard Smith's [P0945R0], it does provide a solution to aliasing anything as a library facility, without running into the problem that [P0945R0] ran into.

Observe:

```cpp
template <template auto Arg>
struct alias /* undefined on purpose */;

template <typename T>
struct alias<T> { using value = T; };

template <auto V>
struct alias<V> { auto value = V; };

template <template <template auto...> typename Arg>
struct alias<Arg> {
  template <template auto... Args> using value = Arg<Args...>;
};
```

We can then use alias when we *know* what it holds:

```cpp
template <template auto Arg>
struct Z {
  using type = alias<Arg>;
  // existing rules work
  using value = typename type::value; // dependent
}; // ok, parses

Z<int> z1; // ok, decltype(z1)::value = int
Z<1> z; // error, alias<1>::value is not a type
```

10.4 Impacts on the function template overloading

The existence of UTPs requires fixes to template overloading; there aren't two ways about it though, it behaves as a universal match that is the worst match in all cases.

```cpp
template<template auto X> const char* kind_name() { return "type"; }
template<template<template auto...> X> const char* kind_name() { return "template"; }
template<auto X> const char* kind_name() { return "value"; }
```

10.5 Bringing CTAD to `make_unique` et. al.

With the introduction of CTAD (constructor template argument deduction) a discrepancy was created which favors using plain new instead of `make_unique` as the latter needs the template arguments of a template class to be spelled out.

With UTPs we can add overloads to `make_unique` which make CTAD work in these situations:

```cpp
auto a = std::tuple(1, true, 'a');        // ok
auto ap = new std::tuple(1, true, 'a');  // ok

// not ok in C++23.
auto up = std::make_unique<std::tuple>(1, true, 'a');

int arr[] = {1, 2, 3};
```

```
auto s = std::span(arr);                // ok
auto sp = new std::span(arr);           // ok

// not implementable in C++23 as span has a NTTP.
auto sup = std::make_unique<std::span>(arr);
```

An overload to implement this would look something like:

```
template<template<template auto...> class C, typename... Ps>
auto make_unique(Ps&&... ps) {
    return unique_ptr<decltype(C(std::forward<Ps>(ps)...))>(new C(std::forward<Ps>(ps)...));
}
```

Note that the `decltype` is required as there is no deduction guide for plain pointers, this is however not a problem specific to the universal template overload.

10.6 Impacts on the specialization of class templates

Universal template arguments enable what appears like overloading of class templates by specializing a unimplemented primary template taking a pack of UTPs.

```
template <template auto...> struct my_container;

template <typename T> struct my_container<T> {
    my_container(T* data, size_t count);
    // A basic implementation
};

template <typename T, typename A> struct my_container<T, A> {
    my_container(T* data, size_t count);
    my_container(T* data, size_t count, const A& alloc);
    // An implementation using an allocator A
};

template <typename T, size_t SZ> struct my_container<T, SZ> {
    my_container(T* data, size_t count);
    // An implementation with an internal storage of SZ bytes
};

template <typename T> my_container(T*, size_t) -> my_container<T>;
template <typename T, typename A> my_container(T*, size_t, const A&) -> my_container<T, A>;
```

As the primary template is not defined there is only a default constructor implicit deduction guide. The explicit deduction guides select the first specialization unless an allocator object is given in which case the second specialization is selected. To select the third specialization the template arguments must be explicitly given.

As default values can only be given for the primary template this is not a perfect emulation of class template overloading, instead additional specializations are needed to emulate use of the defaulted parameters.

10.7 Impact on the partial specialization of NTTP templates

It is a common pattern in C++ to use SFINAE to constrain template parameter types:

```
template <typename T, typename=void>
struct A;

template <template <typename> typename T, typename U>
```

```
struct A<T<U>,std::enable_if_t<std::is_integral_v<U>>>
{
    // Implementation for templates with integral parameter types
};

template <template <typename> typename T, typename U>
struct A<T<U>,std::enable_if_t<std::is_floating_point_v<U>>>
{
    // Implementation for templates with floating point parameter types
};

template <typename T>
struct X {};

A<X<int>> a; // Uses integral partial specialization
```

The covariant behavior of `template auto` allows a similar pattern to be used for NTTPs.

```
template <typename T, typename=void>
struct B;

template <template <template auto> typename T, auto U>
struct B<T<U>,std::enable_if_t<std::is_integral_v<decltype(U)>>>
{
    // Implementation for templates with integral parameter types
};

template <template <template auto> typename T, auto U>
struct B<T<U>,std::enable_if_t<std::is_floating_point_v<decltype(U)>>>
{
    // Implementation for templates with floating point parameter types
};

template <int I>
struct Y {};

B<Y<5>> b; // Uses integral partial specialization
```

This pattern cannot currently be used with NTTPs as `auto` behaves contravariantly:

```
template <typename T, typename=void>
struct C;

template <template <auto> typename T, auto U>
struct C<T<U>,std::enable_if_t<std::is_integral_v<decltype(U)>>>
{};

template <int I>
struct Z1 {};

template <auto I>
struct Z2 {};

C<Z1<5>> c1; // Error: T does not match Z1
C<Z2<5>> c2; // Ok: NTTP can only be `auto`
```

With partial specialization, covariant behavior is desired. We rely on SFINAE to check that `T` accepts `U` and wish to accept any type of template which can take an integral NTTP in this example. `template auto` allows such a pattern to work without modifying the behavior of `auto`.

10.8   Impacts on the specialization of variable templates

UTPs can be used to implement what appears like overloading of variable templates.

The problem of not being able to delete the base case then becomes more pressing than currently as selecting the base case should trigger an error, but will be selected when none of the specialziations matches. This is solved by [P2041R0], which is assumed in the example below.

```cpp
// Metafunction to find a tuple element by a type predicate.
template <template <typename> typename Pred, size_t Pos, typename Tuple>
constexpr size_t tuple_find() {
    if constexpr (Pos == tuple_size_v<Tuple>())
        return npos;
    else if constexpr (Pred<remove_cvref_t<tuple_element_t<Pos, Tuple>>>::value)
        return Pos;
    else
        return tuple_find<Pred, Pos + 1, Tuple>();
}
template <template <typename> typename Pred, typename Tuple>
constexpr size_t tuple_find() { return tuple_find<Pred, 0, Tuple>(); }

// Helper to bind the first arguments of a provided template
template <template <template auto...> typename TPL, template auto... Bs> struct curry {
    template <template auto... Ts> using func = TPL<Bs..., Ts...>;
};

// Unimplemented base case.
template <template auto... Ps>
constexpr size_t tuple_find_v = delete;

template <template <typename> typename Pred, typename Tuple>
constexpr size_t tuple_find_v<Pred, Tuple> = tuple_find<Pred, Tuple>();

template <template <typename> typename Pred, size_t Pos, typename Tuple>
constexpr size_t tuple_find_v<Pred, Pos, Tuple> = tuple_find<Pred, Pos, Tuple>();

// Convenience specialization for use with binary predicate
template <template <typename, typename> typename Pred, typename M, typename Tuple>
constexpr size_t tuple_find_v<Pred, M, Tuple> = tuple_find_v<curry<Pred, M>::template func, Tuple>;

// Convenience specialization to match particular type.
template <typename T, typename Tuple>
constexpr size_t tuple_find_v<T, Tuple> = tuple_find_v<std::is_same, T, Tuple>;
```

This example contains a metafunction `tuple_find` to find a matching element in a tuple based only on its type. Unfortunately it must be implemented as a constexpr function in C++23 if we want it to be usable with or without the start position (which would mimic the `std::find` function in the value domain).

With UTPs we can specialize a template variable `tuple_find_v` to regain the symmetry with current tuple oriented metafunctions such as `tuple_size_v` and `tuple_element_t`.

Given further overloads of the constexpr function `tuple_find` we could simplify the variable template definition

to:

```
template <template auto... Ps> constexpr size_t tuple_find_v = tuple_find<Ps...>();
```

This relies on the power of UTPs in another way, and has the same simplicity as the current variable templates and type aliases of the standard library type traits.

To take the consistency a step further `tuple_find` could be implemented as a class template with UTPs as shown in the previous example instead of as the function it must be in C++23.

11 Choice of keyword (syntax)

The `template auto` syntax in this proposal is a placeholder. EWG needs to decide on a spelling.

This section is written to aid this process.

11.1 Drawbacks of `template auto`

The initially suggested spelling `template auto` is close to existing syntax to explicit function template specialization.

```
template <typename T> auto f(T x) { return x; }

template auto f<float>(float y);
```

11.2 Choice types

We can do any of the following:

— repurpose an existing keyword (like `register`, `inline` or `constexpr`)
— use a combination of keywords (see table)
— make a new keyword
— put a ? after a keyword

A distinct feature like UTPs would generally require a new keyword.

Introducing a new keyword has backwards-compatibility difficulties. Nontheless, C++23 introduced 8 new keywords, including the relatively common words `concept` and `requires`.

This paper proposes a specialist feature, which led us to try to use a compound keyword instead of introducing a new one. *This is novel.*

As shown by the previous section subtle use cases of `template auto` went undetected for years, and even 1.5 implementations were made without detecting it as the grammar productions where the two uses are valid are distinct.

In conclusion introducing a new keyword has a risk of breaking old code, while using two keywords in combination to mean a distinct thing is unprecedented in C++ and carries some risk of its own.

Therefore we suggest an initial poll to make the selection whether to search for a usable and understandable combination of current keywords, or to select a new keyword with suitably obscure spelling to avoid too much code breakage.

To aid in this decision we have brainstormed a set of reasonable token combinations and a set of possible new keywords. None of the lists is exhaustive and additional suggestions are welcomed by the authors.

Keyword combinations	New Keywords
template auto	any_name
auto template	any_kind

Keyword combinations	New Keywords
auto typename	auto_kind
typename auto	unknown_kind
template?	universal_name
	universal_kind
	univ_name
	univ_kind
	indeterminate
	indet_name
	dependent_name
	dep_name

11.2.1 Keyword discussion

If a new keyword is to be recommended the author's list contains two word combinations except for one long and complicated word that could be useable in isolation.

Apart from the collision risk it is of course important that the keyword describes the semantics of the feature as closely as possible. Succinctly the feature semantics can be stated as:

A template parameter that can be bound to any kind of template argument.

This sentence conveys a few facts:

— It is a template parameter
— It can be bound to any kind
— It can have a name (as can all template parameters)

Of these facts we can disregard the first as it is conveyed by the context where the keyword is used, and thus need not be indicated by the keyword. The second fact is the central idea and the word *any* is what conveys this information. The third fact is also given by the context but the spelling `typename` still includes `name` which is to be noted.

This said the list of suggested names include mostly synonyms of `any` in a wide sense: any, auto, unknown, universal, indeterminate, dependent. The word `dependent` was included as it exactly describes how a UTP works to a C++ expert. We also included somewhat obscure abbreviations of universal and dependent which could help reduce the amount of code breakage.

For the last half of the keyword we only came up with name or kind. We think that name is probably best although it may viewed as redundant. One reason is that typename exists and another is that kind seems to refer to a reflection of a UTP, which indicates which *kind* it was bound to. When reflections on UTPs the kind it was bound to for a certain instantiation will be of interest.

11.2.2 To underscore or not to underscore

While the keyword parts are written separately in the table above the intent is to either write the words together or with an interposed underscore.

Checking the current C++ keywords list the keywords consisting of two English words sometimes have an underscore, sometimes not, in a fairly even mix. The authors can't see a pattern which could direct whether to use an underscore in a new keyword. We tested hypotheses that old keywords had no underscore and that longer words requires an underscore, but could not find any correlation.

There are 15 keywords without underscore: alignas, alignof, bitand, bitor, consteval, constexpr, constinit, decltype, inline, noexcept, nullptr, sizeof, typedef, typeid, typename.

There are 12 keywords with underscore: and_eq, co_await, co_return, co_yield, const_cast, dynamic_cast, not_eq, or_eq reinterpret_cast, static_assert, static_cast, thread_local xor_eq.

After considering these aspects the authors came up with the recommendation xxxxxxx which we feel conveys the semantics and is odd enough to not break much code. TODO: Scan github!

11.2.3 Second and third polls

A second poll is either to select a suitable combination of tokens or to select the preferred new keyword spelling recommended from LWG(I?).

If a keyword is selected in the first poll a third poll regarding underscore or not is warranted. It is the opinion of the author that more two word lower case identifiers in C++ code have an underscore than not, therefore making a keyword without underscore less likely to break code.

11.3 Integration with reflection

The authors expect that code using reflection will have a need for this facility. We should not reach for code-generation when templates will do, and so being able to pattern-match on the result of the reflection expression might be a very simple way of going from the `consteval` + injection world back to template matching.

The examples in this section are pending discussion with reflection authors.

Importantly, the authors do not see how one could write `is_specialization_of` with the current reflection facilities, because one would have no way to pattern-match. This is, however, also pending discussion with the authors of reflection.

11.4 Library fundamentals II TS detection idiom

12 Discussion on contravariance correctness

Introduce a way to specify a truly UTP that can bind to anything usable as a template argument.

There are two ways we can go about specifying this, depending on whether we want to be mathematically correct, or less powerful, imprecise, but "easy" to use. (See the co/contravariance discussion below).

12.1 Loose version

We spell it `template auto`. The syntax is somewhat up for debate in this version of the solution.

```
template <template <template auto...> typename F, template auto... Args>
using apply = F<Args...>;

template <typename Type> struct G {};
template <int i> struct H {};
template <template <typename> typename Template> struct J {};

apply<G, int>; // OK, G<int>
apply<H, 3>;   // OK, H<3>
apply<J, G>;   // OK, J<G>
apply<J, H>;   // error, H takes int.
```

12.2 Pedantic version

We need to introduce both an *anything* parameter (which we'll spell `template auto` as above), and a *wildcard* parameter (which, given the work in pattern matching, we are probably free to spell `__`).

We can then define `apply` as follows:

```
template <template <__...> typename F, template auto... Args>
using apply = F<Args...>;
```

`fwd` is similar:

```
template <template <__> typename F, temmplate auto arg>
using fwd = F<arg>;
```

In this mechanism, `__` really is an "unconstrained" pattern match, and does not declare a parameter kind - it's illegal to declare a `template <__ x> struct foo {};`, for instance.

12.2.1 Example of the difference between mathematically correct and easy versions.

The basic premise about "mathematically correct" is that the user of the template template parameter must be able to instantiate it in all ways that the declaration indicates. This means that if a template parameter is declared `template<template auto...> typename F` only an argument with the same declaration can be substituted, as this is the only way that F can be instantiated with any number of template parameters of any kind.

In the "Easy" version it is enough that the template argument can be instantiated with *some* template parameters out of the ones permitted by the template parameter declaration.

```
template<template<template auto...> class TPL> void f()
{
    TPL<int, 3> myObject;        // #1
}


f<std::vector>();               // #2
f<std::array>();                // #3
```

With the easy definition both instantiations at #2 and #3 are fine, but then we get substitution failure for #2 at #1 as `std::vector` can't take a numerical second template argument. At #3 we get no errors as f() *happened* to instantiate TPL only with template argument kinds that the `std::array` supports.

With the mathematically correct definition both #2 and #3 fail to substitute into TPL as none of `std::array` and `std::vector` have a variadic UTP, and thus can't support whatever instances that f() may want to create.

To be able to pass any template as the argument to a template parameter in the mathematically correct definition we must use another syntax for the template parameters of the template template parameter, such as ___. This allows any class template to be used as the template argument and defers checking until inside the instantation process.

```
template<template<__...> class TPL> void f()
{
    TPL<int, 3> myObject;        // #1
}


f<std::vector>();               // #2
f<std::array>();                // #3
```

Here #2 fails late at #1 and #3 succeeds just as for the easy definition.

13 Covariance and Contravariance

13.1 As of C++23

Consider the two concepts:

```
template <typename T> concept A = true;
// B subsumes A
template <typename T> concept B = A<T> && true;
```

Covariance	Contravariance

```
auto returns_a() -> A;              template <template <B> typename f>
auto returns_b() -> B;              using puts_b = void;
auto f() {                          template <template <A> typename f>
    // OK, requirement less constrained than re using puts_a = void;
    A x = returns_b();              template <B> using takes_b = void;
    // Error, requirement stricter than returne template <A> using takes_a = void;
    B y = returns_a();              using x = puts_b<takes_a>; // ok
}                                   // Error, constraint mismatch (gcc)
                                    // clang accepts (in error)
                                    using w = puts_a<takes_b>;
```

We are used to the covariant case, but the usage of contravariant cases is not as common. The issue with `puts_a<takes_b>` is that `puts_a` requires a metafunction with a *wider interface* than one that just accepts Bs.

We have seen that concept-constrained template parameters behave correctly - covariantly on returns, contravariantly on parameters; but do other template parts as well?

Let's just replace A with `auto` and B with `int`:

```
template <template <int> typename f> using puts_int = void;
template <template <auto> typename f> using puts_auto = void;
template <int> using takes_int = void;
template <auto> using takes_auto = void;
using x = puts_int<takes_auto>; // OK
using w = puts_auto<takes_int>; // Error, but MSVC, GCC and clang all accept
```

TODO: Ask James Touton where in the standard this is made an error.

Function pointers do not convert in either co- or contravariant ways:

```
struct X {}; struct Y : X {};
using f_of_x = void(*)(X&);
using f_of_y = void(*)(Y&);
// Error, no conversions between function pointers
f_of_y fy = static_cast<f_of_x>(nullptr);

using f_to_x = X&(*)();
using f_to_y = Y&(*)();
// Error, no conversions between function pointers
f_to_x xf = static_cast<f_to_y>(nullptr);
```

It does work for virtual function covariant return types:

```
struct ZZ {};
struct Z : ZZ {
    virtual auto f() -> Z&;
    virtual void g(Z&);
};
struct W : Z {
    auto f() -> W& override; // OK, covariant return type
    // Error, no contravariant parameter types
    void g(ZZ&) override;
};
```

How about parameter packs?

```
template <template <typename> typename f> using puts_one = void;
template <template <typename...> typename f> using puts_var = void;
template <typename> using takes_one = void;
template <typename...> using takes_var = void;
using x = puts_one<takes_var>; // OK in 17 and 20, Error in 14
using w = puts_var<takes_one>; // OK, compiles
```

Turns out parameter packs behave both ways (but also covariantly) - not what one would expect given the concepts example above.

13.2 Design space exploration for `template auto`

We have inconsistent behavior across the language. Let's say we did the right thing and made `template auto` behave contravariantly, like concepts.

```
template <template <auto> typename f> using puts_value = void;
template <template <template auto> typename f> using puts_any = void;
template <auto> using takes_value = void;
template <template auto> using takes_any = void;
using x = puts_value<takes_any>; // OK
using w = puts_any<takes_value>; // Error because contravariant.
```

But then, how do we write `apply`? Let's do it for a single argument to avoid complications with a covariant . . . :

```
template <template <template auto T> typename f, template auto arg>
using apply1 = f<arg>;
```

The above is correct, but *useless* - it requires `f`s signature to be `template <template auto T>`. What we want to express is that `f` is any kind of unary template metafunction, so we can pass in something like `template <int x> using int_constant = std::integral_constant<int, x>;`.

As a thought experiment, let's call the covariant version of `template auto` (with the meaning "deduce this from the argument") `__`:

```
template <template <__ T> typename f, template auto arg>
using apply1 = f<arg>;
```

This is what we want to express (and check at instantiation time), but now the `arg`s constraint is spelled differently from `f`'s constraint, and that might be very, very difficult to teach.

It also requires us to reserve an additional combination of tokens. Contrast what happens if we just made `template auto` behave covariantly (the way `__` behaves above) if used in that position. What do we lose?

At first glance, we lose the ability to define the contravariant meaning of `template auto` - a template parameter that *can* bind to anything. But can we get that back?

Recall that concepts behave contravariantly. Consider this one:

```
namespace std {
  template <template auto arg>
  concept anything = true;
};
```

If we relaxed the rule that only type concepts could appear in the template parameter list, we could say this:

```
template <template <std::anything ARG> typename takes_anything,
          template auto arg>
using apply1 = takes_anything<arg>;
```

This relaxation is *unlikely to happen* (given the past oral arguments in EWG), but we can still constrain with `requires`, albeit that is a far inferior-looking strategy.

This would behave *contravariantly*! While `template auto ARG` means "deduce", a metafunction taking a concept (which behaves contravariantly!) that *anything* satisfies *has* to be a metafunction taking `template auto` (or `std::anything`).

In light of this, the paper authors have come to the conclusion that trying to make `template auto` behave contravariantly is all downside and little upside. The example to follow with `template auto` should be the behavior of `...` (deduction behavior, both co- and contravariant).

This allows for usage to look the same as declaration, and like to bind to like. We anticipate the feature being much easier to teach this way, and in the rare cases when someone really needs the contravariant behavior, they will know to use a library-provided concept. It is also consistent with the rest of the non-concept template language.

Of course, the library-concept solution requires an additional relaxation of constraints syntax, but at least it's feasible in the future, as opposed to requiring a second difficult-to-teach token sequence.

14   Other Considered Syntaxes

In addition to the syntax presented in the paper, we have considered the following syntax options:

14.1   `.` and `...` instead of `template auto` and `template auto ...`

```
template <template <...> typename F, . x, . y, . z>
using apply3 = F<x, y, z>;
```

The reason we discarded this one is that it is very terse for something that should not be commonly used, and as such uses up valuable real-estate.

15   Compendium: Open Design Questions

See discussions above to inform these choices; this section just compiles the design questions for EWG and EWGI.

15.1   Eager or Late Checking

Whether to:

— specify **eager** checking with conservative rules on use of identifiers, or;
— specify **late** checking on instantiation and use the C++23 rules for parsing dependent expressions for them.

15.2   Easy or Correct

Whether to:

— (**correct**) reserve *two* token sequences (`template auto` and `__`) and have a single defined meaning for both, or;
— (**easy**) reserve *one* token sequence (`template auto`) and switch its meaning in a context-dependent manner, losing the ability to mean `template auto` in a *template template-parameter*, since in that context, `template auto` would mean `__`.

16 Acknowledgements

17 References

[P0522R0] James Touton, Hubert Tong. 2016-11-11. DR: Matching of template template-arguments excludes compatible templates.
https://wg21.link/p0522r0

[P0634R3] Nina Ranns, Daveed Vandevoorde. 2018-03-14. Down with typename!
https://wg21.link/p0634r3

[P0945R0] Richard Smith. 2018-02-10. Generalizing alias declarations.
https://wg21.link/p0945r0

[P1830R1] Ruslan Arutyunyan. 2019-10-07. std::dependent_false.
https://wg21.link/p1830r1

[P1936R0] Ruslan Arutyunyan. 2019-10-07. Dependent Static Assertion.
https://wg21.link/p1936r0

[P2008R0] Mateusz Pusz. 2020-01-10. Enable variable template template parameters.
https://wg21.link/p2008r0

[P2041R0] David Stone. 2020-01-11. Deleting variable templates.
https://wg21.link/p2041r0

[P2098R0] Walter E Brown, Bob Steagall. 2020-02-17. Proposing std::is_specialization_of.
https://wg21.link/p2098r0

[P2593R0] Barry Revzin. 2022-05-21. Allowing static_assert(false).
https://wg21.link/p2593r0

[P2601R1] Justin Cooke. 2022-07-16. Make redundant empty angle brackets optional.
https://wg21.link/p2601r1